



■ Torch no-overlap ■ Triton-Dist ■ AG old / RS new ■ AG new / RS old ■ AG new / RS new


Attention: A100 (4 GPUs)




Attention: K100_AI (4 GPUs)




MLP: A100 (4 GPUs)




MLP: K100_AI (4 GPUs)



Full E2E: A100 (4 GPUs)



Full E2E: K100_AI (4 GPUs)



Each platform is normalized to its own Torch no-overlap record; no cross-platform absolute comparison is implied.